

VENKAT.CHUNDURU
venkatchunduru@outlook.com
408-569-3138

Technical Profile:

- Hands on big data **professional with proven track record** of developing and delivering large scale impactful projects in large enterprises and startups with Hadoop Spark NoSQL etc.,
- Expertise in building near real time actionable insights analytics and reporting platforms with effective **data ingestion and transformation frameworks**, innovation in creating enterprise solutions suitable to next gen use cases.
- **Architected and designed**, Big Data applications using Scala, Java, Spark Structured Streaming, Akka, Spark SQL, Hive, MongoDB, Impala, HBase, Phoenix, Flume, Kafka, Oozie, Sqoop, Solr, Parquet, Avro, CDH parcels and csd etc.,
- Expertise with containers and container orchestration services like Kubernetes, marathon, dockers etc., and microservices containerization
- Developed reliable fault tolerant self-healing applications capable of processing near real time and batch data sets, also applications for SaaS PaaS IaaS
- Familiarity with development tools like IntelliJ, Eclipse etc., hands on exposure developing in Java Scala, Frameworks like **Spring Boot , Camel REST Akka** etc., and build tools like maven Ant sbt Gradle etc.,
- Proficiency with SQL in NoSQL and RDBMS like MySQL Oracle RDS MariaDB etc.,
- Diligent programmer and quick learner with **excellent problem solving**, presentation and mentoring skills.
- Strong communication and interpersonal skills.

Unique Skills:

- Customer first approach in supporting and implementing product features
- Performance tuning, identifying bottlenecks and provide with creative solutions with minimum to zero impact to business
- Laser focus on timeline and quality of deliverable via coordinating with team and setting expectations
- Design and architect with keen attention to resource utilization, and creating a lean and efficient applications
- Always keep technology stack to current or latest stable versions.
- Come up with alternate solutions in critical scenarios and production issues.
- Large volume high throughput platforms and frameworks.
- Explore and use cutting edge tools solution to support critical business use cases.

Professional Experience:

Oracle America Inc. Redwood Shores, CA

August 2021 – Present

cx marketing - Senior Principal Data Engineer

As member of cx marketing a division of Oracle Marketing Cloud, I have been involved in customer loyalty and rewards management group (Crowdtwist) and events flow processing team that manages marketing campaign performance (Responsys). As part of Crowdtwist data team responsible for on time processing of transactional data to Redshift data mart. This data is accessed via Yellowfin to get different reports on the effectiveness of campaigns. As part of Responsys developed Infinity data collector service application that integrates both infinity and responsys events.

Responsibilities:

- Developed a new application that would process the marketing campaign events generated from mobile app via Infinity service using Web Sockets
- This application would allow Responsys platform to integrated events from mobile app and email campaign.
- Application is created using Akka typed actors for processing and Akka streams for infinity services integration via Web Sockets
- Support current cloverDX job flows and make performance improvements, make changes to logic that determines delta data.
- Created an architecture and design to move from current cloverDX ETL processing to spark EMR cluster (Oracle's Data Flow) based processors.
- Worked resolving the production issues, identified root cause of stuck ETL processes that cause delays up to days to catchup. Implemented preventive measures.

Environment : Oracle Cloud (OCI) , Kubernetes, Helm charts, Akka typed actors , Akka streams, WebSocket, Amazon cloud, Redshift ,Oracle, Aurora DB, CloverDX, cloudWatch, python, spark, spark sql, Yellowfin, Java , eclipse.

Collective[i]. San Jose, CA

May 2019 – March 2021

(Cross Commerce media, New York)

Director – Bigdata Engineering

Collective[i] merges client CRM and other relevant sources of data into our network to generate intelligence used by the world's leading sales organization to generate revenue, streamline/automate non-revenue producing activities, and enable enterprises to be customer-centric and data-driven. Our patent-pending technology and applications designed for sales managers, operators and professionals are essential to the modern sales force.

Responsibilities:

- As a director of big data engineering managed entire data processing of application.
- **Improved performance** of data processing there by reducing the processing time to half leading to most current data available to customers.

- Upgraded the current cloudera hadoop from 5.11 to 6.2 version and upgraded all the current application to version match latest version of cloudera.
- Upgraded current spark applications from 1.6/2.0 to spark 2.4 and scala 2.11
- Architected and implemented a master data management process that ensure all the data is associated to same master data across the application
- Manage upgrade and performance tune current ETL job which are based on Hive Impala, udfs, spark, phoenix scheduled via oozie coordinators for oozie workflows
- Upgraded current HBase Phoenix based application to current version of apache phoenix and recreated with all the required indexes
- Improved the performance and **reliability** of Impala queries as the UI application has queries that directly serve data via impala queries.
- Redesigned the current Architecture of Mongo Hive processors to improve the performance and reliability of the application.
- Redesigned the current Architecture of Hive Neo4j processor to improve the performance and reliability of the connectors builder
- Architected designed and implemented Intelligence write back component of the product, that keeps CRM and Intelligence applications in sync.
- Created **reliable SLA** for ADAS (Analytics Data as Service) application that interact with frontend application intelligence.com and Hadoop data lake.
- Evaluated and created a strategy to move the current on prem application to cloud environment with minimal impact to current application
- Created a migration path from on prem to cloud migration with lift and shift and operationalize to a modernized architecture.
- Evaluated the data lake engine Dermio for the purpose of adhoc querying of the data lake.

Environment: Cloudera Hadoop (CDH 5.11 / 6.2), Spark 2.0/ 2.4 (Structured Streaming), Kafka, oozie , Akka, Scala 2.11, MongoDB, Hive, Beeline, Impala, Neo4j 4.2.2, Java 1.8, Phoenix 5.0.0, Bitbucket, HBase, zookeeper, HUE, HDFS, yarn, Parquet, Python, Jenkins, Redis, Elastic search

Verizon, Palo Alto

Principal Engineer – Platform Engineering

Sep 2015 – May 2019

As part of Verizon labs team, the project is to create a re-usable bigdata platform which consists of **re-usable** ingestion, anonymize, processing, enriching frameworks suitable to create actionable insights, analytics and reports. Ability to handle different kind of data sets like **click stream** or network or **IOT devices** data.

Responsibilities:

- Architected designed and implemented a reusable **dynamically configurable** rest based akka http based data ingestion services suitable for any kind of data sets. Is deployed in production supporting video analytics data and IOT devices data.
- Added ability to access processing data schema from external sources like Couchbase or Kafka to data process framework (ETL framework).
- Added spark structured streaming capability to ETL framework
- Designed and implemented streaming services that can accept and process and ingest to HDFS via Kafka which would be consumed by ETL pipeline.

- Updated an existing ETL pipeline from spark hive impala to spark Kafka druid pipeline.
- Designed and implemented a http client utility that can be used to stream data to Hadoop data cluster with ability to provide a **seamless authentication and configurable anonymization**.
- Integrated ingestion from different data sources like click stream, IOT thingspace, sensity, visp etc.,
- Supporting batch ingestion pipeline that consume data that landed on to a GFS cluster processed via spark process that processes and anonymizes suitable for consumption by ETL pipeline

Environment: Cloudera Hadoop (CDH 5.7.1), Spark 2.1 (Structured Streaming), Kafka, Druid, Akka, Scala 2.11, Couchbase, Hive, Beeline, Impala, Maven, CentOS, MySql, Java 1.8, Solr 6.5.1, Git, Stash, HBase, zookeeper, HDFS, yarn, Avro, Parquet, Python, Apache Camel, JMeter, Jenkins, Kubernetes, Mesos, Marathon etc.,

Apple Inc., Cupertino, ca

Oct 2013 – Sep 2015

Senior Consultant – Big Data initiative
(MQM R&D)

As part of Strategic Data Solutions (SDS) 's MQM R&D team, the project is to set up Hadoop cluster to capture manufacturing data generated by different manufacturing sites across the world, which would be used to prevent fraud and abuse of apple products in addition to monitoring manufacturing quality and yield in real time.

CA Technologies, Sunnyvale, ca
(Arcotsystem Inc.,)

Jan 2013 - Oct 2013

Solutions Consultant

Serene Corporation, Santa Clara, CA

March'2012- Jan'2013

Principal Consultant

Cisco, San Jose, CA

May'2010-March'2012

Sr. Architect

YouSendIt, Campbell, CA

May'2009-May'2010

Staff Engineer

Symantec Inc., Mountain View, CA

April'2008-May'2009

Senior Software Engineer/Consultant

Palm, Inc. Sunnyvale, CA

Sep'2007-April'2008

Senior Software Engineer/Consultant

Apple Inc., Cupertino, CA

Jan'2007- Sep'2007

Senior Software Engineer/Consultant

VeriSign Inc., Mountain View, CA

May'2006- Jan'2007

Senior Software Engineer/Consultant

Vodafone Inc., Walnut Creek, CA

Sep'2005-May'2006

Lead Engineer

SBC Inc., San Ramon, CA

Feb'2005-Sep'2005

Lead Engineer

Harris Broadcasting Corporation, Sunnyvale, CA

June'2003-Feb'2005

Lead Engineer

Education

- Bachelor of Engineering Bangalore University, India.